


# Huxley-Gödel Machine

Human-Level Coding Agent  
Development by an  
Approximation of the Optimal  
Self-Improving Machine

KAUST, 2025



# The Challenge



## Central Question

Which self-modifications  
should we accept?

### ✗ Prior Approach

DGM & SICA use benchmark scores  
Assume: High score → Better lineage

### ⚠ The Problem

High scores ≠ Good descendants


# Metaproductivity-Performance Mismatch

## Performance

Immediate scores


## Metaproductivity

Long-term potential



# Visualizing the Mismatch

## ✗ High Score, Poor Lineage



Trend: ⬇ Declining

CMP ≈ 0.70 (max descendant)


## ✓ Lower Score, Great Lineage



Trend: ⬆ Improving

CMP ≈ 0.95 (max descendant)

# Self-Improvement as Tree-Search



At each step:

**Modify** agent → create child

**Evaluate** agent on task

**Goal:** Find best final agent

# From Gödel Machine to HGM



**Gödel Machine**

Formal proofs  
Theoretically optimal




**HGM**

Estimate CMP  
Practical approximation

# Clade-Metaproductivity (CMP)

The Idea: Focus on **Lineages**



**Clade** = Agent + all descendants

Why CMP?

- ✓ Modest ancestor can have great descendants
- ✓ More statistically robust
- ✓ Inspired by biology (Huxley)
- ✓ Captures long-term potential

# Theorem 1: CMP Oracle = Gödel Machine

**CMP oracle is sufficient to implement the Gödel Machine**



CMP  $\equiv$  Q-value



Provably optimal



No proofs needed



**Insight:** Estimate CMP  $\rightarrow$  Approximate Gödel Machine

# Deep Dive: What is a Gödel Machine?

**Self-referential universal problem solver** that makes provably optimal self-improvements

## Core Mechanism

Runs a **proof searcher** looking for formal proofs that a code modification will increase expected utility

## When proof found

The modification is **executed** and permanently alters the machine

## Guarantee

**Provably optimal** with respect to given objective

## Challenge

Finding proofs is **computationally expensive** and impractical



Schmidhuber (2003)

# Assumption 1: Our Specialized Setting

How does our **coding agent development** setting differ from general Gödel Machine?

## 1 Final-only objective

Only care about the **final agent**, not intermediate rewards

## 3 No proof cost

Proofs don't consume budget (unlike original GM)

## 2 Repeatable trials

Can **reset environment** for each test; evaluations are independent

## 4 Uniform costs


Each self-modification costs **exactly 1 budget unit**



These simplifications make optimal policy **tractable** via CMP!

# Global Metaproductivity (GMP)


How does evolving agent  $a$  affect the **entire tree's** final performance?



⚠️ **Global Impact**  
Expanding  $a_0$  affects ALL future nodes

$$\text{GMP}_{\pi}(\mathcal{T}, a) = \mathbb{E} \left[ U \left( \max_{a' \in \mathcal{T}_B} a' \right) \right]$$

Expected utility of best agent in ENTIRE final tree



⚡ **Key insight:**

GMP = Q-value in RL

(state = tree, action = expand agent)

# From GMP to CMP: Why Clades?

## ✗ Problem with GMP

- Too global - considers entire tree
- Self-modifications can affect ancestors
- Hard to conceptualize
- Difficult to estimate

## ✓ Solution: CMP

- Localized to **subtree** (clade)
- Only descendants matter
- Conceptually clear
- Practically estimable

## ▣ Gödel Machine Focus


Decides whether to **accept or reject** based on provable potential of **subsequent** self-improvements

## 🌿 Biological Inspiration

**Clade** (Huxley, 1957): lineages of common ancestry

# GMP vs CMP: Visual Comparison


**GMP: Global Scope**



**Considers ALL nodes**

Max over  $\mathcal{T}_B$  (entire tree)

**CMP: Clade Scope**



**Only descendants of  $a_1$**

Max over  $C(\mathcal{T}_B, a_1)$  (clade of  $a_1$ )

# The Mathematics of CMP

Clade-Metaproductivity: Expected utility of best agent in **clade**

$$\begin{aligned}\text{CMP}_\pi(\mathcal{T}, a) &= \mathbb{E}_{\mathcal{T}_B \sim p_\pi(\cdot | \mathcal{T}, a)} \left[ U \left( \arg \max_{a' \in C(\mathcal{T}_B, a)} \text{Score}_\pi(a') \right) \right] \\ &= \mathbb{E}_{\mathcal{T}_B \sim p_\pi(\cdot | \mathcal{T}, a)} \left[ \max_{a' \in C(\mathcal{T}_B, a)} U(a') \right] \quad (\text{if Score} = U)\end{aligned}$$

where  $C(\mathcal{T}_B, a)$  is the clade (subtree rooted at  $a$ ) in final tree  $\mathcal{T}_B$

 **Estimator**

 **Where**


 **Weighted**

$$n_{\text{success}}^C(a) = \sum_{a' \in C(a)} n_{\text{success}}(a')$$

More evaluated agents get higher weight

# CMP Estimator: Step-by-Step

How to estimate CMP: Aggregate test results across the **clade**



## Step 1: Collect in Clade

Gather all results from  $a_1$  and descendants:

$a_1: 7✓, 3✗ | a_2: 5✓, 2✗ | a_3: 4✓, 6✗ | a_4: 8✓, 1✗ | a_5: 6✓, 4✗$

## Step 2: Sum Totals

$$n_{\text{success}}^C(a_1) = 7 + 5 + 4 + 8 + 6 = 30$$

$$n_{\text{failure}}^C(a_1) = 3 + 2 + 6 + 1 + 4 = 16$$

## Step 3: Compute CMP

$$\widehat{\text{CMP}}(a_1) = \frac{30}{30 + 16} = \frac{30}{46} \approx 0.652$$

# Theorem 1: The Main Result

## Theorem 1

Under **Assumption 1**, access to the **CMP oracle** is sufficient to implement the **Gödel Machine**.

### 🔑 Key Equality

$$\text{CMP}_\pi(\mathcal{T}, a) = Q_\pi(\mathcal{T}, a)$$

CMP is exactly the Q-value in the Gödel POMDP!

### 🎯 POMDP Setup

- **State:** Tree + parent + child
- **Action:** Accept or reject
- **Observe:** Only parent, child, budget

### ✓ Optimality

Selecting max CMP = selecting max Q-value = **Bellman optimal**



**Implication:** No need for expensive proofs! Just estimate CMP well.

# Proof Sketch: Why CMP = Q-value

## 1 Gödel Machine Policy

At each step: observe  $(a_{\text{parent}}, a_{\text{child}}, b)$  and choose to accept or reject child

## 2 Final Selection

$\text{Score}_\pi$  is indicator function: picks either final parent or final child

## 3 Clade Structure

Only descendants reachable; ancestors above clade are **not accessible**

## 4 Expected Return


$$\begin{aligned}\text{CMP}_\pi((\mathcal{T}, a_p, a_c, b), a) &= \mathbb{E}[U(\arg \max_{a' \in C} \text{Score}(a'))] \\ &= \mathbb{E}[U(\arg \max_{a' \in \{a_p, a_c\}} \text{Score}(a'))] \\ &= Q_\pi((\mathcal{T}, a_p, a_c, b), a)\end{aligned}$$

## ✓ Conclusion

CMP oracle knows the true Q-values → can make optimal accept/reject decisions

# POMDP: The Gödel Machine's Decision Process

Gödel Machine operates as a POMDP: Partial observability, Accept/Reject decisions




## Transition

1. Generate child via self-modification
2. New parent = accepted agent
3. Budget decreases by 1

## Reward

Final reward =  $U(\text{Score}_\pi)$  picks best from final (parent, child) pair

# Why CMP = Q-value: Visual Proof



**Clade of a\_child**  
Only these nodes are reachable!

## 1 Final Selection Constraint

Score $_{\pi}$  is an **indicator function**: Picks either  $a_{\text{parent}}$  OR  $a_{\text{child}}$ . All other agents get score 0

## 2 Clade = Reachable Set

If we accept  $a_{\text{child}}$ , only its **descendants** are reachable.  $C(\mathcal{T}_B, a_{\text{child}}) =$  all nodes below  $a_{\text{child}}$

## 3 The Equality

$$\text{CMP} = \mathbb{E}[\max_{a' \in C} U(a')] = \mathbb{E}[\max_{a' \in \{a_p, a_c\}} U(a')] = Q_{\pi}$$

# Thompson Sampling for Exploration-Exploitation

How to select which agent to expand when we only have **estimates** of CMP?



## Thompson Sampling

1. Model CMP as Beta distribution:

$$\text{Beta}(\tau(1 + n_s^C), \tau(1 + n_f^C))$$

2. Sample from each posterior

3. Select agent with highest sample



## Exploration Scheduler

Temperature parameter  $\tau(t)$  increases over time:

- **Early:** Low  $\tau \rightarrow$  more exploration
- **Late:** High  $\tau \rightarrow$  more exploitation


Automatically balances discovery vs refinement



Unlike greedy baselines, TS **probabilistically** explores promising clades

# Temperature Schedule: $\tau(t)$ Over Time

Temperature  $\tau(t) = \frac{B}{b}$  increases over time: more **exploitation** as budget runs out



## Early (low $\tau$ )

Wide Beta distributions → more random sampling → **exploration**

## Middle (medium $\tau$ )

Moderate distributions → balanced → **mixed strategy**

## Late (high $\tau$ )

Narrow distributions → pick best → **exploitation**

## Beta Distributions Evolution

### Low $\tau = 1$

$$\text{Beta}(1 + 7, 1 + 3)$$

Wide curve

### Med $\tau = 5$

$$\text{Beta}(5 + 35, 5 + 15)$$

Medium

### High $\tau = 20$

$$\text{Beta}(20 + 140, 20 + 60)$$


Peaked

## Effect on Selection

- **Low  $\tau$ :** Sample values spread out → try different agents
- **High  $\tau$ :** Sample values concentrated → pick highest CMP

# Connection to Infinite-Armed Bandits

When to **expand** (create new agent) vs **evaluate** (test existing)?



## 📊 UCB-Air Strategy

Expand when:

$$N_t^\alpha \geq |\mathcal{T}_t|$$

Number of evaluations vs number of agents

## ⚡ Decoupled Policy

- DGM/SICA: Always expand then evaluate
- HGM: Adaptive choice each step

# HGM: How to Estimate CMP?

Aggregate success across **entire clade** (lineage)

$$\widehat{\text{CMP}}(a) = \frac{\text{successes in clade}}{\text{total tests in clade}}$$



Clade-level  
aggregation



Weighted by  
evaluations



Thompson  
Sampling

# HGM Algorithm: Complete Walkthrough

Step-by-step execution of one HGM iteration

## Step 1: Compute CMP Estimates

For each agent  $a$  in tree:

$$\widehat{\text{CMP}}(a) = \frac{n_{\text{success}}^C(a)}{n_{\text{success}}^C(a) + n_{\text{failure}}^C(a)}$$

## Step 2: Thompson Sample

For each  $a$ , sample:

$$S_C(a) \sim \text{Beta}(\tau(1 + n_s^C), \tau(1 + n_f^C))$$

Select  $a^* = \arg \max_a S_C(a)$

## Step 3: Decide Action

If  $N_t^\alpha \geq |\mathcal{T}_t|$ : **Expand** (create child)

## Step 4a: If Expanding

1. Agent  $a^*$  self-modifies → creates child  $c$
2. Add  $c$  to tree  $\mathcal{T}$
3. Initialize counters for  $c$

## Step 4b: If Evaluating

1. Sample agent to test (by individual stats)
2. Run on one task → get result
3. Update  $n_{\text{success}}$  or  $n_{\text{failure}}$
4. Bubble up to ancestors (update  $n^C$ )

## Step 5: Repeat

Continue until budget exhausted. Return best-belief agent:

# HGM: Three Adaptive Policies



## Expansion

Which agent to modify?

Use **clade CMP**



## Evaluation

Which agent to test?

Use **individual stats**



## Selection

Expand or evaluate?

**Adaptive** scheduling



**Key Innovation:** Decoupled expansion from evaluation!

# HGM in Practice



## **Asynchronous**

Run on all CPUs  
Early stopping



## **Best-Belief**

Select final agent  
by posterior

# Experimental Setup



## Benchmarks

SWE-bench Verified (500)  
SWE-bench Lite (300)  
Polyglot



## Baselines


DGM  
SICA  
SWE-agent



## Models

GPT-5 / GPT-5-mini  
Qwen3-Coder

# Main Results from Paper



# Result 1: CMP Correlation

| Method                     | Weighted Correlation | Unweighted Correlation |
|----------------------------|----------------------|------------------------|
| SICA                       | 0.444 / 0.274        | 0.444 / 0.274          |
| DGM                        | 0.285 / 0.383        | 0.406 / 0.357          |
| <b>HGM</b>                 | <b>0.778 / 0.626</b> | <b>0.512 / 0.873</b>   |
| SWE-Verified-60 / Polyglot |                      |                        |



HGM: **2-3x better correlation** with true metaproductivity

# Result 2: Performance & Efficiency

## 🎯 Accuracy

SWE-Verified-60

**56.7%**

+16.7% improvement

## ⚡ Speed

vs DGM

**2.38×**

faster

Polyglot

**30.5%**

+10.2% improvement

Polyglot speedup

**6.86×**

faster

# Result 3: vs. Human Designers

**HGM on SWE-bench Verified**

**53.2%**

Initial agent

**61.4%**

HGM discovered 



**Top-10** on leaderboard (all models)  
**#1** GPT-5-mini based system

# Result 4: Human-Level Performance!

HGM optimized with **GPT-5-mini** → Evaluated with **GPT-5**

| Agent                         | SWE-Lite Standard (%)   |
|-------------------------------|---|
| SWE-agent (Best human design) | 56.7  |
| <b>HGM + GPT-5</b>            | <b>57.0</b>  |

 Transfers across  
**model sizes**

 Optimized on  
**different dataset**

 Not overfitting,  
**genuine ability**

# Emergent Behaviors



## Self-Motivated Iteration

Agents perform **multiple self-modifications** per instruction

→ *Arbitrary levels of meta-improvement!*



## Nested Diff Structures

Diff patches of diff patches  
Multiple levels of changes

*"Mind-bending to understand manually"*

# Key Contributions



**Identified** Metaproductivity-Performance Mismatch



**Introduced** Clade-Metaproductivity (CMP)



**Proved** Theorem 1 ( $CMP = \text{Gödel Machine}$ )



**Developed** HGM Algorithm



**Validated** 2-3 $\times$  better CMP estimation



**Achieved** Human-level performance

# HGM vs. Baselines

|                    | SICA        | DGM           | HGM  |
|--------------------|-------------|---------------|--|
| <b>Guidance</b>    | Performance | Performance   |  <b>CMP</b> |
| <b>Expansion</b>   | Greedy      | Probabilistic | <b>Thompson Sampling</b>   |
| <b>Decoupled?</b>  | ✗           | ✗             | ✓  |
| <b>Theory</b>      | ✗           | ✗             | ✓ <b>Gödel Machine</b>   |
| <b>Correlation</b> | 0.44        | 0.29          | <b>0.78</b>  |
| <b>Speed</b>       | 1×          | 1×            | <b>2-7×</b>  |

# Key Takeaways



**Performance ≠ Metaproductivity**



**Lineages > Individuals**



**CMP Oracle = Gödel Machine**



**Human-Level Performance**



**Paradigm Shift:** Focus on capacity to **keep improving**